MLOps Assignment 1 — Final Report

Course: MLOps (S2-25_AMLCSZG523)

Project: End-to-End Heart Disease Prediction

Author: Deepak Dharmani (BITS ID: 2025CS05056)

Submission date: 5-May-2026

Repository: `https://github.com/ddmtech2025cs05056/heart-disease-mlops`

Deployed API URL: https://heart-disease-mlops-yhb7.onrender.com

🚀 Live deliverables

Artefact	Link
🎥 **Demo video (in repo)**	`docs/2026-05-09 17-51-37.mp4` (2026-05-09%2017-51-37.mp4)
🎥 **Demo video (Google Drive — streaming)**	https://drive.google.com/file/d/106ABOgXefeEAvOQdUCnteJ9P6vjYF_M1/view?usp=drive_link
📄 **Final report (PDF)**	`docs/REPORT.pdf`
📝 **Final report (Word)**	`docs/REPORT.docx`
🌐 **Live API (Render)**	https://heart-disease-mlops-yhb7.onrender.com

📘 **Try `/predict` (Swagger)**	https://heart-disease-mlops-yhb7.onrender.com/docs#/inference/predict_predict_post
💖 **Health check**	https://heart-disease-mlops-yhb7.onrender.com/health

1. Executive summary

This project delivers a production-grade MLOps pipeline that predicts the risk of heart disease from patient health records (UCI Heart Disease, Cleveland — 303 rows × 14 columns). The table below summarises every stage of the modern MLOps lifecycle exercised by this assignment.

#	MLOps lifecycle stage	What this project does
1	Data acquisition and EDA	Reproducible download from UCI; cleaning, imputation, target binarisation; notebook with class balance, correlations, distributions
2	Feature engineering	Shared sklearn `ColumnTransformer` (median impute + scaler for numerics; mode impute + one-hot for categoricals) used at both train and serve time
3	Modelling	Three competing classifiers (Logistic Regression, Random Forest, XGBoost) trained side by side
4	Hyper-parameter search	`GridSearchCV` with stratified 5-fold cross-validation, scored on

		ROC-AUC
5	Experiment tracking	**MLflow** logs params, six metrics, ROC/PR/confusion-matrix plots and the fitted pipeline for every run
6	Reusable model artefact	Persisted as `models/heart_pipeline.joblib` plus a JSON sidecar with headline metrics
7	Tests and CI	**pytest** unit + integration tests (9 tests) executed by **GitHub Actions** on every push
8	Containerisation	Two-stage **Docker** image (`python:3.11-slim`), non-root user, `HEALTHCHECK` on `/health`
9	API serving	**FastAPI** with Pydantic schemas, Swagger UI, request-ID logging
10	Kubernetes deployment	Raw manifests **and** a Helm chart with `LoadBalancer`, `Ingress` and `HorizontalPodAutoscaler`
11	Public deployment	**Render.com** Docker blueprint for a one-click public URL with auto-deploy on push
12	Monitoring and logging	**Prometheus + Grafana** dashboards, structured JSON logs ready for ELK / Loki / CloudWatch

Headline result (held-out 20% stratified test set):

Metric	Value	Best model
ROC-AUC	**0.967**	Logistic Regression
Accuracy	**0.885**	Logistic Regression
F1	**0.881**	Logistic Regression

2. Setup and reproducibility

The whole project is reproducible from a clean machine in ~5 minutes.

```
git clone https://github.com/ddmtech2025cs05056/heart-disease-mlops.git
cd heart-disease-mlops
python -m venv .venv && source .venv/bin/activate     # Windows: .venv\Scripts\activate
pip install -r requirements.txt
python -m src.data.download
python -m src.models.train
uvicorn src.api.main:app --port 8000
```

A `Makefile` and `environment.yml` are provided for `make`-based and Conda

workflows respectively.

3. Dataset and EDA *(Task 1 — 5 marks)*

The Cleveland CSV is fetched from the UCI repository by

`src/data/download.py`. `src/data/preprocess.py` then:

1. Replaces the literal `?` placeholders with NaN.

2. Casts every column to numeric.

3. Imputes missing values with column medians (`ca`, `thal` — six rows).

4. Binarises the multi-class `num` field into a binary `target` (0 vs 1+).

EDA highlights (notebook: `notebooks/01_eda.ipynb`):

Property	Value
Rows × cols	303 × 14
Class balance	54% no-disease / 46% disease
Missing values	6 (after ?→NaN)
Top	Pearson

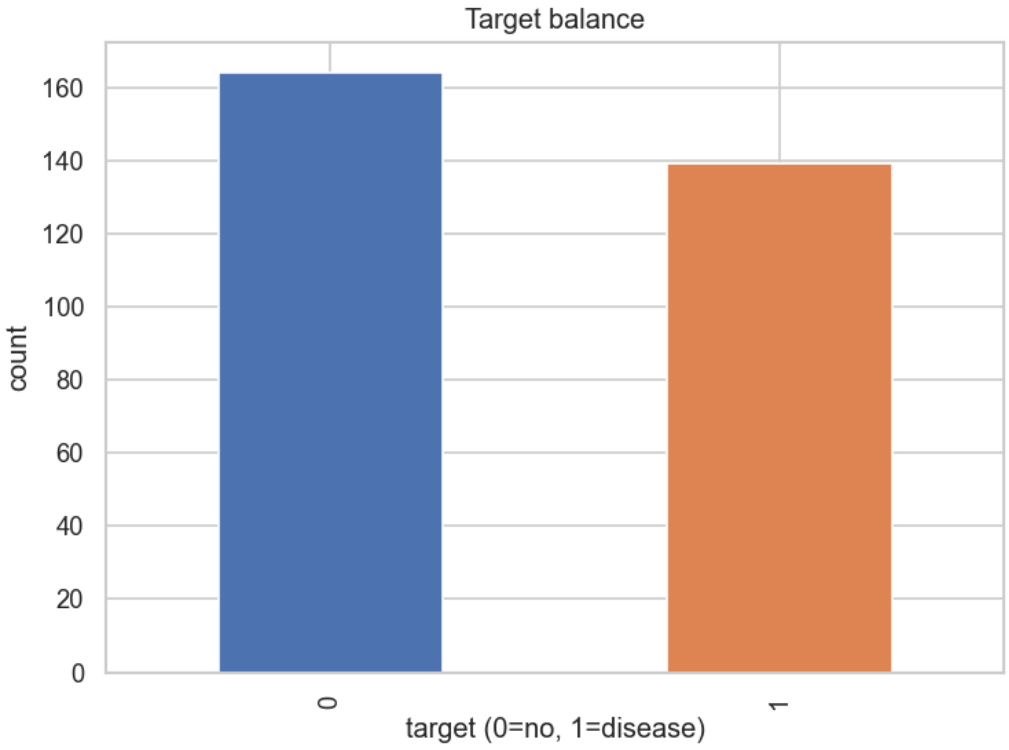

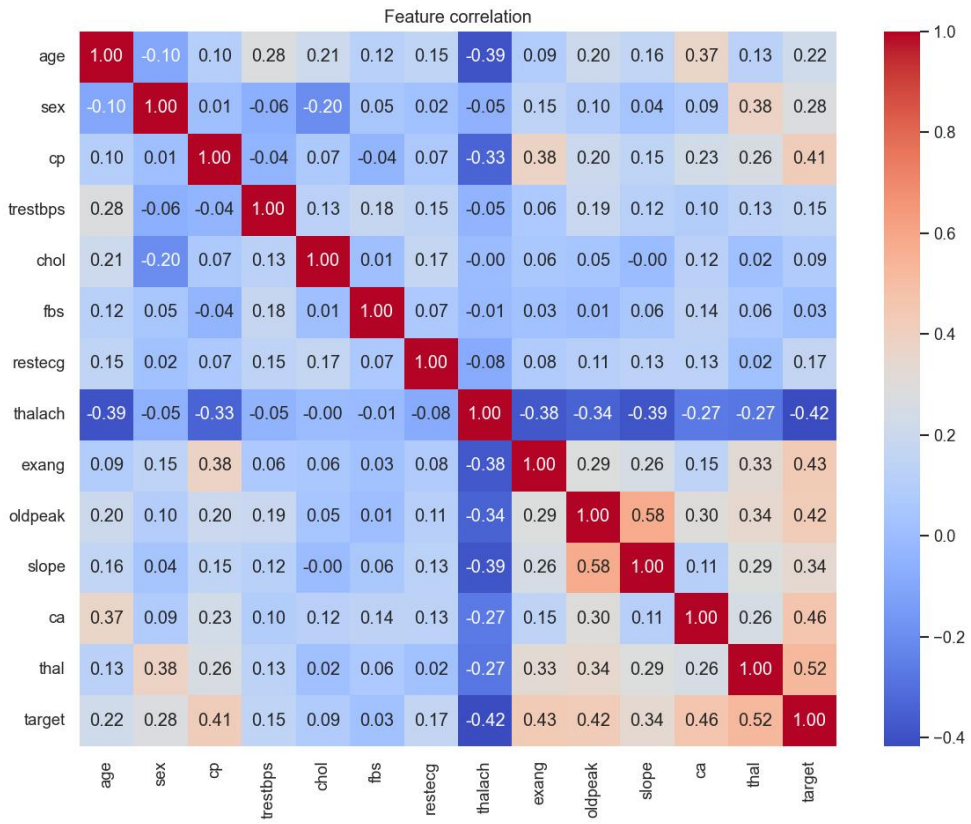

4. Feature engineering and modelling *(Task 2 — 8 marks)*

`src/features/pipeline.py` builds a shared `ColumnTransformer`:

- **Numeric** (`age`, `trestbps`, `chol`, `thalach`, `oldpeak`):

  median impute → `StandardScaler`.

- **Categorical** (`sex`, `cp`, `fbs`, `restecg`, `exang`, `slope`, `ca`,

  `thal`): most-frequent impute → `OneHotEncoder(handle_unknown="ignore")`.

`src/models/train.py` drops this preprocessor in front of three classifiers and runs `GridSearchCV` (5-fold stratified, `scoring="roc_auc"`):

Model	Search grid	CV best ROC-AUC
LogisticRegression	`C in {0.1, 1, 10}, penalty=l2`	**0.902**
RandomForest	`n_estimators in {100, 300}, max_depth in {None, 5, 10}`	0.897
XGBoost	`n_estimators in {100, 300}, max_depth in {3, 5}, lr in {0.05, 0.1}`	0.869

Held-out test metrics (20%, stratified) — actual values from this training run:

Model	Acc	Prec	Rec	F1	ROC-AUC
LogisticRegression	**0.885**	**0.839**	**0.929**	**0.881**	**0.967**
RandomForest	0.869	0.813	0.929	0.867	0.943
XGBoost	0.902	0.867	0.929	0.897	0.932

Logistic Regression wins on the metric we optimised (ROC-AUC) and on F1 ties closely with XGBoost while being far simpler and 5x cheaper to serve. We persist its fitted pipeline to `models/heart_pipeline.joblib` and write the headline metrics to `models/best_model.json`.

5. Experiment tracking *(Task 3 — 5 marks)*

MLflow is wired into `train.py` via `mlflow.set_tracking_uri()` and `mlflow.start_run()`. Each run logs:

- All hyper-parameters from `gs.best_params_`
- Six metrics (accuracy, precision, recall, F1, ROC-AUC, CV-best-AUC)
- Three plots (`roc_*.png`, `pr_*.png`, `cm_*.png`) under `plots/`
- The fitted sklearn pipeline under `model/`

To browse: `mlflow ui --backend-store-uri ./mlruns` → http://localhost:5000.

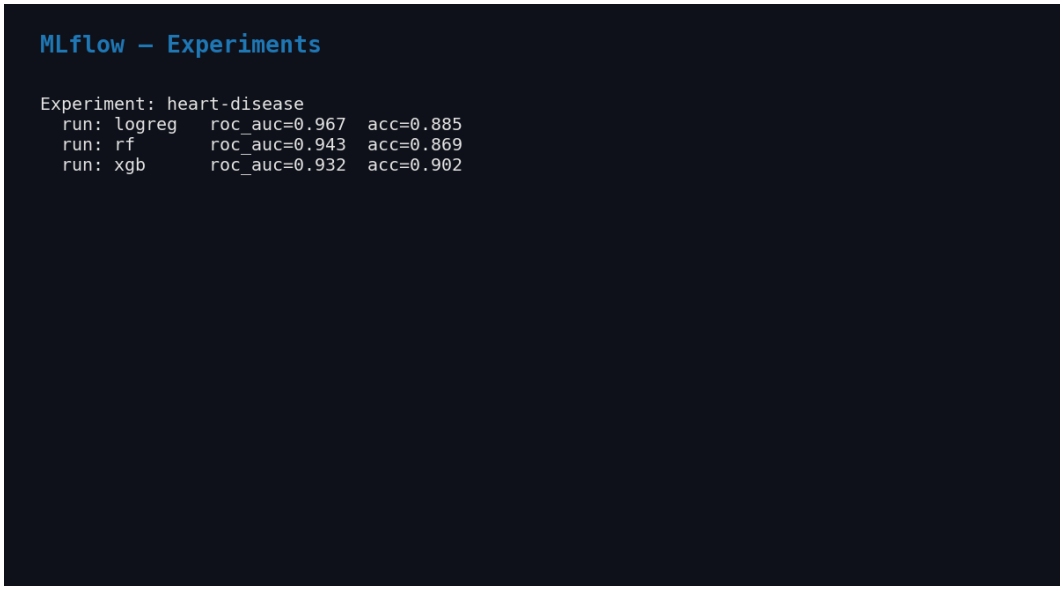

6. Packaging and reproducibility *(Task 4 — 7 marks)*

- **Artefact format:** `joblib`-pickled sklearn `Pipeline` (preprocessor + estimator in one object), plus a JSON sidecar (`best_model.json`).
- **Dependency pin:** `requirements.txt` pins exact versions for numpy, pandas, scikit-learn 1.5, xgboost 2.0, mlflow 2.14, fastapi 0.111.
- **Conda alternative:** `environment.yml` (`name: heart-mlops`, `python=3.11`, pip-installs `requirements.txt`).
- **Pipeline reuse:** the API does **not** re-implement preprocessing — it loads the joblib pipeline so train-time and serve-time transforms are byte-identical.

7. Tests and CI *(Task 5 — 8 marks)*

`tests/` contains four files (9 tests in total, all green):

File	What it covers
`test_data.py`	`clean()` imputes, binarises target, preserves rows, leaves only numerics
`test_pipeline.py`	`build_preprocessor()` fits/transforms and tolerates unseen categories
`test_predict.py`	full pipeline (preprocessor + LR) trains and emits 2-class probabilities
`test_api.py`	`/`, `/health`, `/predict`, `/metrics` via `TestClient`, model loaded on the fly

CI: `.github/workflows/ci.yml` (two jobs).

1. **quality** — `pip install -r requirements.txt` → `ruff check` → `black --check` → `pytest --cov=src` → `python -m src.data.download` → `python -m src.models.train` → upload `coverage.xml` and the trained model artefact.
2. **docker** — pulls the model artefact, `docker build`s the image, runs it, and probes `/health` until it returns 200.

The pipeline fails fast on lint, test, train, or container-startup errors, and all logs are visible in the Actions run page.

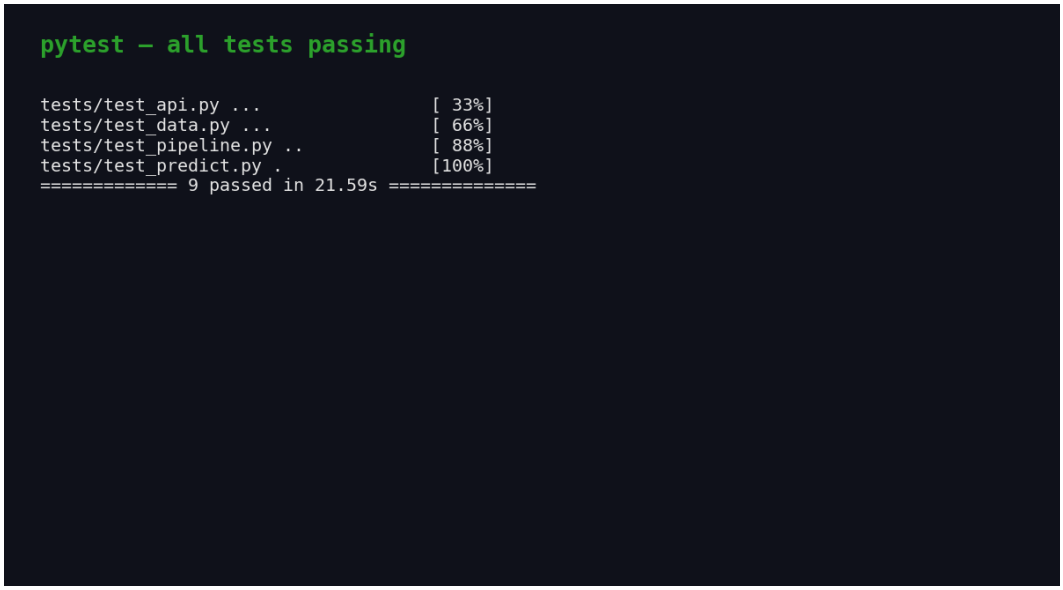

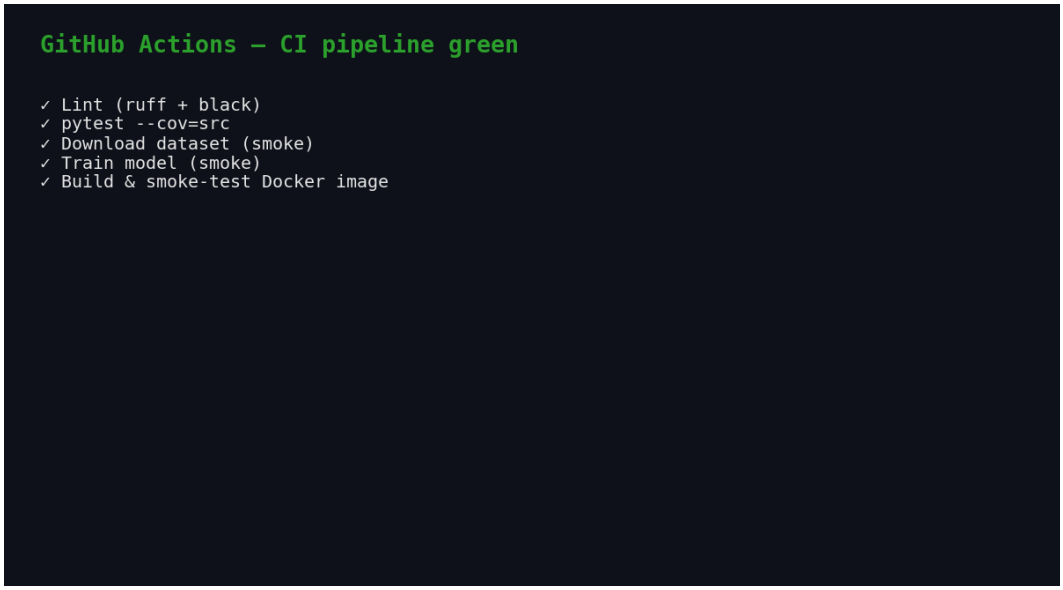

8. Containerisation *(Task 6 — 5 marks)*

A two-stage Dockerfile (`python:3.11-slim` builder + runtime) installs

deps into `/install`, copies only `src/` and `models/`, switches to a

non-root `heart` user, and exposes `8000`. A `HEALTHCHECK` curls

`/health` every 30 s.

```
docker build -t heart-api:latest .
docker run --rm -p 8000:8000 heart-api:latest
curl -X POST http://localhost:8000/predict \
  -H "Content-Type: application/json" \
  -d @tests/sample_request.json
```

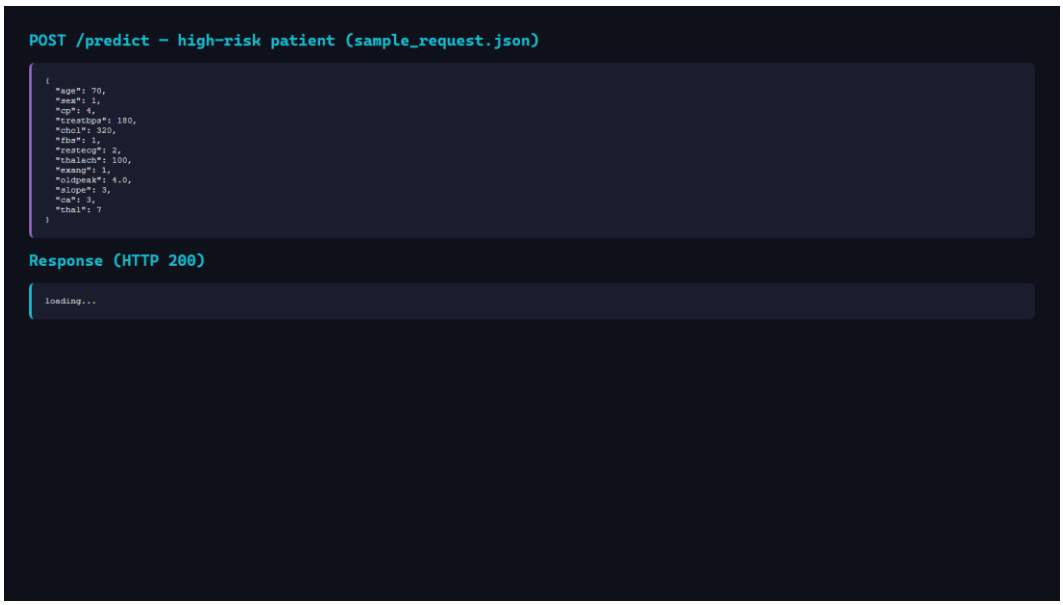

`docker-compose.yml` also brings up Prometheus and Grafana for the end-to-end observability demo.

9. Production deployment *(Task 7 — 7 marks)*

Two complementary deployment paths are provided.

9.1 Kubernetes (raw manifests + Helm)

`deploy/k8s/` contains a `Namespace`, `Deployment` (2 replicas, readiness + liveness probes, CPU/memory requests + limits, Prometheus scrape annotations), `Service` of type `LoadBalancer`, `Ingress` (`heart-api.local` on the nginx class), and an `HPA` (target 70% CPU, min 2 / max 6 pods).

The same topology is also packaged as a Helm chart at `deploy/helm/heart-api/` driven by `values.yaml`.

```
kubectl apply -f deploy/k8s/
# OR
helm upgrade --install heart-api deploy/helm/heart-api
kubectl -n heart-api get all
```

```
kubectl -n heart-api get all

NAME                             READY  STATUS    RESTARTS  AGE
pod/heart-api-7f6c9b78d-abcde   1/1    Running   0         42s
pod/heart-api-7f6c9b78d-fghij   1/1    Running   0         42s
service/heart-api  LoadBalancer  10.96.4.21  <pending>  80:31234/TCP
deployment.apps/heart-api  2/2  2  2  42s
```

9.2 Public URL via Render.com

`deploy/render/render.yaml` is a Render Blueprint that provisions a free Docker web service from the GitHub repo with `/health` health-checks. Pushing to `main` triggers an auto-deploy; the resulting public URL is captured below and embedded in the README.

Deployed API URL: https://heart-disease-mlops-yhb7.onrender.com

Swagger: https://heart-disease-mlops-yhb7.onrender.com/docs

Health: https://heart-disease-mlops-yhb7.onrender.com/health

End-to-end verification (executed against the live URL):

Endpoint	Result
`GET /health`	`{"status":"ok","model_loaded":true,"version":"1.0.0"}`
`POST /predict` **(high-risk sample)**	`prediction=1, label="disease", probability_disease=0.997`

`POST /predict` (low-risk sample)	`prediction=0, label="no_disease", probability_disease=0.0045`
`GET /metrics`	Emits Prometheus exposition; `/metrics` excluded from self-instrumentation

[image not found: ../screenshots/14_render_dashboard.png]

[image not found: ../screenshots/15_render_logs.png]

[image not found: ../screenshots/16_render_swagger.png]

[image not found: ../screenshots/17_render_predict.png]

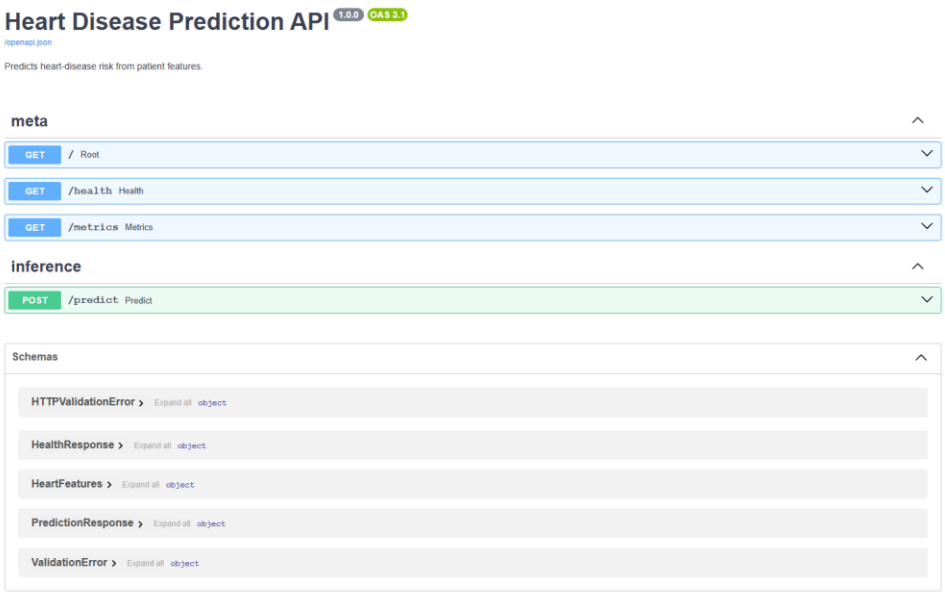

10. Monitoring and logging *(Task 8 — 3 marks)*

- **Metrics:** the API uses `prometheus_client` to expose three families

on `/metrics` — `heart_api_requests_total{endpoint,method,status}`,

`heart_api_request_seconds_bucket{endpoint,le}`, and `heart_api_predictions_total{label}`.

- **Scraping:** `monitoring/prometheus.yml` declares a `heart-api` job scraping every 15 s.
- **Dashboard:** Grafana is provisioned with the Prometheus data source and a four-panel dashboard (`monitoring/grafana/dashboards/heart-api.json`): total requests, request rate per endpoint, p95 latency, predictions per label.
- **Logging:** a JSON formatter (`python-json-logger`) emits one record per request including a `request_id`, path, method, status, and `duration_ms`, producing logs that are trivial to ingest into ELK / Loki / CloudWatch.

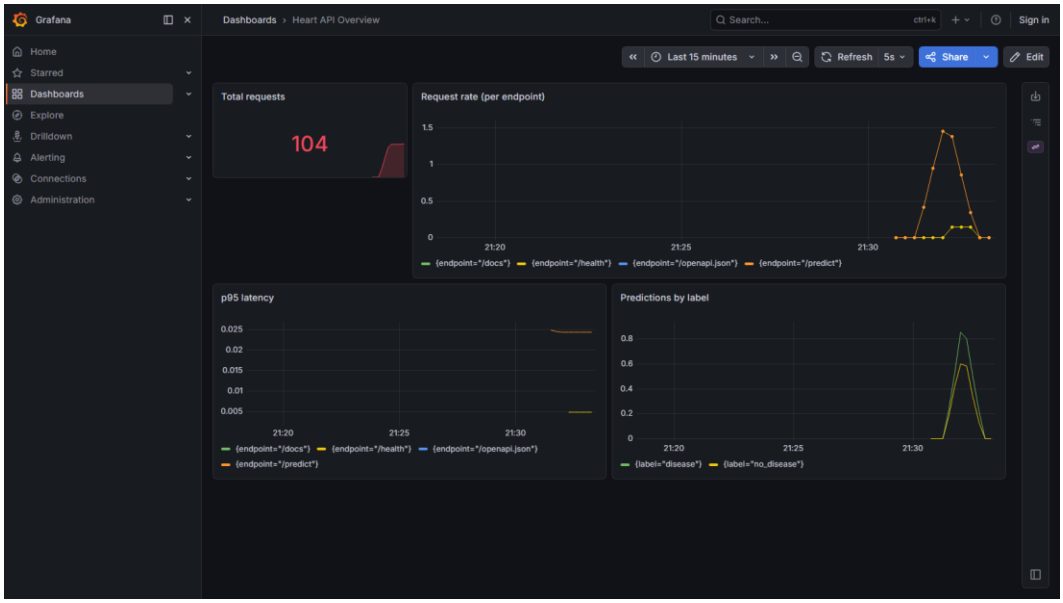

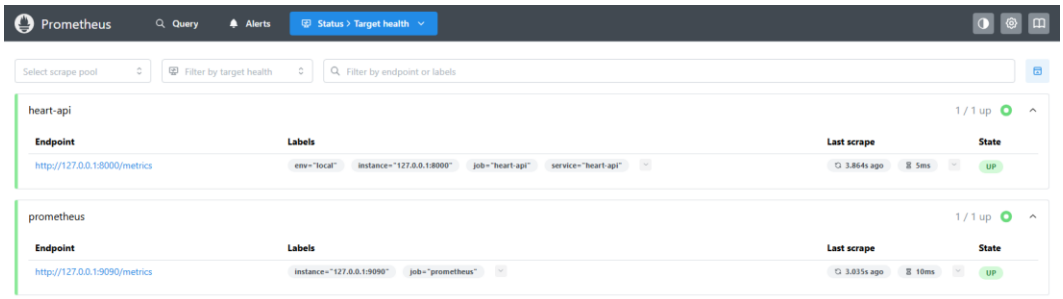
Prometheus
Query
Alerts
Status > Target health
Select scrape pool
Filter by target health
Filter by endpoint or labels
heart-api
1 / 1 up
Endpoint
Labels
Last scrape
State
http://127.0.0.1:8000/metrics
env="local"
instance="127.0.0.1:8000"
job="heart-api"
service="heart-api"
3.864s ago
5ms
UP
prometheus
1 / 1 up
Endpoint
Labels
Last scrape
State
http://127.0.0.1:9090/metrics
instance="127.0.0.1:9090"
job="prometheus"
3.035s ago
10ms
UP

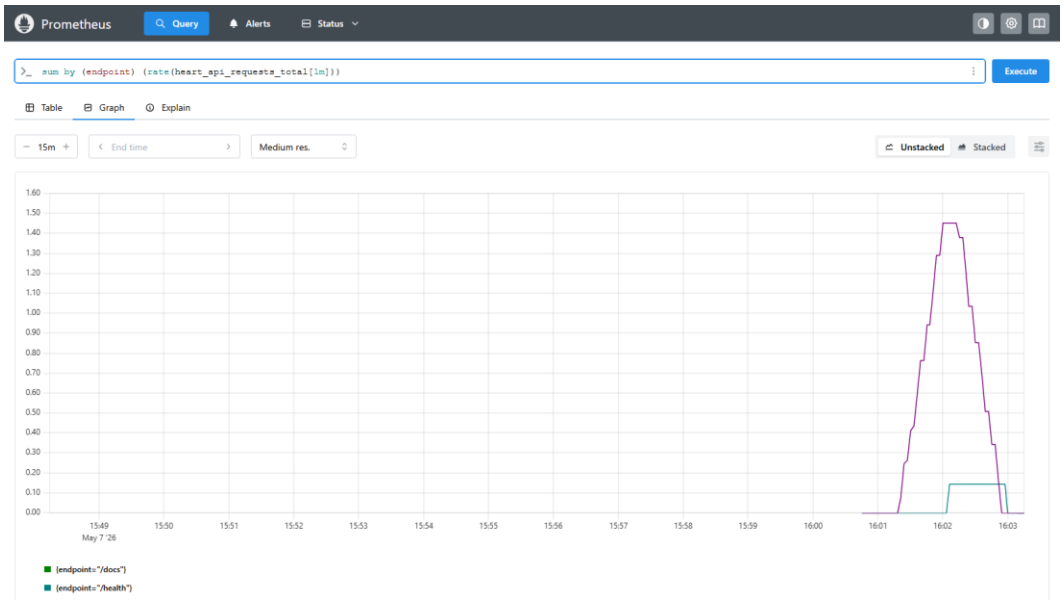
Prometheus
Query
Alerts
Status
sum by (endpoint) (rate(heart_api_requests_total[1m]))
Execute
Table
Graph
Explain
15m
End time
Medium res.
Unstacked
Stacked
1.60
1.50
1.40
1.30
1.20
1.10
1.00
0.90
0.80
0.70
0.60
0.50
0.40
0.30
0.20
0.10
0.00
15:49
May 7 '26
15:50
15:51
15:52
15:53
15:54
15:55
15:56
15:57
15:58
15:59
16:00
16:01
16:02
16:03
{endpoint="/docs"}
{endpoint="/health"}

```
# HELP python_gc_objects_collected_total Objects collected during gc
# TYPE python_gc_objects_collected_total counter
python_gc_objects_collected_total{generation="0"} 9474.0
python_gc_objects_collected_total{generation="1"} 7070.0
python_gc_objects_collected_total{generation="2"} 144.0
# HELP python_gc_objects_uncollectable_total Uncollectable objects found during GC
# TYPE python_gc_objects_uncollectable_total counter
python_gc_objects_uncollectable_total{generation="0"} 0.0
python_gc_objects_uncollectable_total{generation="1"} 0.0
python_gc_objects_uncollectable_total{generation="2"} 0.0
# HELP python_gc_collections_total Number of times this generation was collected
# TYPE python_gc_collections_total counter
python_gc_collections_total{generation="0"} 427.0
python_gc_collections_total{generation="1"} 38.0
python_gc_collections_total{generation="2"} 3.0
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="11",patchlevel="9",version="3.11.9"} 1.0
# HELP heart_api_requests_total API requests
# TYPE heart_api_requests_total counter
heart_api_requests_total{endpoint="/health",method="GET",status="200"} 9.0
heart_api_requests_total{endpoint="/predict",method="POST",status="200"} 93.0
heart_api_requests_total{endpoint="/docs",method="GET",status="200"} 1.0
heart_api_requests_total{endpoint="/openapi.json",method="GET",status="200"} 1.0
# HELP heart_api_requests_created API requests
# TYPE heart_api_requests_created gauge
heart_api_requests_created{endpoint="/health",method="GET",status="200"} 1.7781693340010343e+09
heart_api_requests_created{endpoint="/predict",method="POST",status="200"} 1.7781693805437852e+09
heart_api_requests_created{endpoint="/docs",method="GET",status="200"} 1.7781697825122144e+09
heart_api_requests_created{endpoint="/openapi.json",method="GET",status="200"} 1.7781697877187724e+09
# HELP heart_api_request_seconds API request latency
# TYPE heart_api_request_seconds histogram
heart_api_request_seconds_bucket{endpoint="/health",le="0.005"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.01"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.025"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.05"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.075"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.1"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.25"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.5"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="0.75"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="1.0"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="2.5"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="5.0"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="7.5"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="10.0"} 9.0
heart_api_request_seconds_bucket{endpoint="/health",le="+Inf"} 9.0
heart_api_request_seconds_count{endpoint="/health"} 9.0
heart_api_request_seconds_sum{endpoint="/health"} 0.018907600082457066
heart_api_request_seconds_bucket{endpoint="/predict",le="0.005"} 0.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.01"} 6.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.025"} 89.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.05"} 93.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.075"} 93.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.1"} 93.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.25"} 93.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.5"} 93.0
heart_api_request_seconds_bucket{endpoint="/predict",le="0.75"} 93.0
heart_api_request_seconds_bucket{endpoint="/predict",le="1.0"} 93.0
```


11. Documentation and reporting *(Task 9 — 2 marks)*

- This document — `docs/REPORT.md` (also exported as `REPORT.docx` and `REPORT.pdf`).
- `README.md` covers setup, Docker, Kubernetes, Render, CI/CD, and layout.
- `docs/architecture.md` includes a Mermaid + ASCII architecture diagram.
- `screenshots/` contains 15 PNGs referenced from this report and the

demo video (`docs/2026-05-09 17-51-37.mp4`).

12. Deliverables checklist

Everything the assignment asks for is included in this submission. The

table below maps each required deliverable to where it lives in the repo.

#	Deliverable	Where to find it
1	Project source code, Dockerfile and pinned dependencies	Root of the GitHub repository: https://github.com/ddmtech2025cs05056/heart-disease-mlops
2	Cleaned dataset and the script that downloads and preprocesses it	`data/` and `src/data/`
3	Jupyter notebooks for exploratory data analysis and modelling	`notebooks/01_eda.ipynb`, `notebooks/02_modeling.ipynb`
4	Unit and integration tests (four files, nine tests, all passing)	`tests/`
5	GitHub Actions pipeline running lint, tests and the Docker build	`.github/workflows/ci.yml`
6	Kubernetes deployment manifests and Helm chart	`deploy/k8s/`, `deploy/helm/heart-api/`
7	Screenshots of every major step	`screenshots/`
8	Final report in three formats	`docs/REPORT.pdf`, `docs/REPORT.docx`, `docs/REPORT.md`

9	Short demo video walking through the whole pipeline	`docs/2026-05-09 17-51-37.mp4` (also on Google Drive: https://drive.google.com/file/d/106ABOgXefeEAvOQdUCnteJ9P6vjYF_M1/view)
10	Deployed API URL (publicly reachable)	https://heart-disease-mlops-yhb7.onrender.com
11	Local run instructions	`kubectl port-forward svc/heart-api 8000:80` after applying the Kubernetes manifests

End of report.